

# NoSQL Data Stores

Ashish Kedia  
Final Year, IT

January 29, 2016

# Outline

- 1 Introduction
- 2 Rise of NoSQL
- 3 Improve RDBMS
- 4 CAP Theorem
- 5 Consistency
- 6 Pros and Cons
- 7 Architecture
- 8 Basic Terminologies
- 9 Basic Operation
- 10 Other Databases
- 11 Tips for using NoSQL

# Introduction

- Stands for No-SQL or **N**ot **O**nly SQL

# Introduction

- Stands for No-SQL or **N**ot **O**nly SQL
- Class of non-relational data storage systems

# Introduction

- Stands for No-SQL or **N**ot **O**nly SQL
- Class of non-relational data storage systems
- Examples include MongoDB, Cassandra, Google Datastore, DynamoDB

# Introduction

- Stands for No-SQL or **N**ot **O**nly SQL
- Class of non-relational data storage systems
- Examples include MongoDB, Cassandra, Google Datastore, DynamoDB
- Do not require a fixed table schema

# Introduction

- Stands for No-SQL or **N**ot **O**nly SQL
- Class of non-relational data storage systems
- Examples include MongoDB, Cassandra, Google Datastore, DynamoDB
- Do not require a fixed table schema
- Distributed data storage systems

# Introduction

- Stands for No-SQL or **N**ot **O**nly SQL
- Class of non-relational data storage systems
- Examples include MongoDB, Cassandra, Google Datastore, DynamoDB
- Do not require a fixed table schema
- Distributed data storage systems
- BASE instead of ACID



# Historical Needs

- Large Dataset. Explosion of storage needs

# Historical Needs

- Large Dataset. Explosion of storage needs
- Dynamically-typed data with frequent schema changes

# Historical Needs

- Large Dataset. Explosion of storage needs
- Dynamically-typed data with frequent schema changes
- High frequency of read-write operations

# Historical Needs

- Large Dataset. Explosion of storage needs
- Dynamically-typed data with frequent schema changes
- High frequency of read-write operations
- Horizontal Scalability

# Historical Needs

- Large Dataset. Explosion of storage needs
- Dynamically-typed data with frequent schema changes
- High frequency of read-write operations
- Horizontal Scalability
- Rise of Cloud Computing

# Historical Needs

- Large Dataset. Explosion of storage needs
- Dynamically-typed data with frequent schema changes
- High frequency of read-write operations
- Horizontal Scalability
- Rise of Cloud Computing
- Spikes in Web Applications

# Improvements in RDBMS

- Caching Mechanism - But limited scalability

# Improvements in RDBMS

- Caching Mechanism - But limited scalability
- Parallel Databases - OLTP unavailable



# Improvements in RDBMS

- Caching Mechanism - But limited scalability
- Parallel Databases - OLTP unavailable
- Write to master, read from slave. Replaced by Paxos

# Improvements in RDBMS

- Caching Mechanism - But limited scalability
- Parallel Databases - OLTP unavailable
- Write to master, read from slave. Replaced by Paxos
- Sharding - Loss of references

# Improvements in RDBMS

- Caching Mechanism - But limited scalability
- Parallel Databases - OLTP unavailable
- Write to master, read from slave. Replaced by Paxos
- Sharding - Loss of references
- Application needs to be Partition Aware

# Improvements in RDBMS

- Caching Mechanism - But limited scalability
- Parallel Databases - OLTP unavailable
- Write to master, read from slave. Replaced by Paxos
- Sharding - Loss of references
- Application needs to be Partition Aware
- Multi-Master, In-Memory, Insert Only, No Join

# Brewer's CAP Theorem

- consistency, availability and partitions

# Brewer's CAP Theorem

- consistency, availability and partitions
- Availability - 99.999%. Resilient. Large Node System

# Brewer's CAP Theorem

- consistency, availability and partitions
- Availability - 99.999%. Resilient. Large Node System
- choose 2 out of 3 for any shared-data system

# Brewer's CAP Theorem

- consistency, availability and partitions
- Availability - 99.999%. Resilient. Large Node System
- choose 2 out of 3 for any shared-data system
- to scale, we need partition. So choose A or C



# Brewer's CAP Theorem

- consistency, availability and partitions
- Availability - 99.999%. Resilient. Large Node System
- choose 2 out of 3 for any shared-data system
- to scale, we need partition. So choose A or C
- AP Type - Google Datastore, DynamoDB

# Brewer's CAP Theorem

- consistency, availability and partitions
- Availability - 99.999%. Resilient. Large Node System
- choose 2 out of 3 for any shared-data system
- to scale, we need partition. So choose A or C
- AP Type - Google Datastore, DynamoDB
- CP Type - Couchbase, Cassandra

# Consistency in NoSQL

- determines rules for visibility and apparent order of updates

# Consistency in NoSQL

- determines rules for visibility and apparent order of updates
- eventually consistent model

# Consistency in NoSQL

- determines rules for visibility and apparent order of updates
- eventually consistent model
- eventually all updates will propagate through the system

# Consistency in NoSQL

- determines rules for visibility and apparent order of updates
- eventually consistent model
- eventually all updates will propagate through the system
- BASE as opposed to ACID

# Consistency in NoSQL

- determines rules for visibility and apparent order of updates
- eventually consistent model
- eventually all updates will propagate through the system
- BASE as opposed to ACID
- Key-Look ups are always consistent

# ACID

- **Atomic** All of the work in a transaction committed or none of it



# ACID

- **Atomic** All of the work in a transaction committed or none of it
- **Consistent** Consistency is defined in terms of constraints. Always obeyed.

# ACID

- **Atomic** All of the work in a transaction committed or none of it
- **Consistent** Consistency is defined in terms of constraints. Always obeyed.
- **Isolated** Results are not visible until the transaction has committed.

# ACID

- **Atomic** All of the work in a transaction committed or none of it
- **Consistent** Consistency is defined in terms of constraints. Always obeyed.
- **Isolated** Results are not visible until the transaction has committed.
- **Durable** The results of a committed transaction survive failures

# BASE

- Acronym contrived to be the opposite of ACID

# BASE

- Acronym contrived to be the opposite of ACID
- **B**asically **A**vailable,

# BASE

- Acronym contrived to be the opposite of ACID
- **B**asically **A**vailable,
- Soft state,

# BASE

- Acronym contrived to be the opposite of ACID
- **B**asically **A**vailable,
- Soft state,
- Eventually Consistent

# BASE

- Acronym contrived to be the opposite of ACID
- **B**asically **A**vailable,
- Soft state,
- Eventually Consistent
- Weak consistency stale data OK



# BASE

- Acronym contrived to be the opposite of ACID
- **B**asically **A**vailable,
- Soft state,
- Eventually Consistent
- Weak consistency   stale data OK
- Availability first

# BASE

- Acronym contrived to be the opposite of ACID
- **B**asically **A**vailable,
- Soft state,
- Eventually Consistent
- Weak consistency stale data OK
- Availability first
- Best effort. Optimistic Concurrency Control

# BASE

- Acronym contrived to be the opposite of ACID
- **B**asically **A**vailable,
- Soft state,
- Eventually Consistent
- Weak consistency stale data OK
- Availability first
- Best effort. Optimistic Concurrency Control
- Approximate answers OK

# BASE

- Acronym contrived to be the opposite of ACID
- **B**asically **A**vailable,
- Soft state,
- Eventually Consistent
- Weak consistency stale data OK
- Availability first
- Best effort. Optimistic Concurrency Control
- Approximate answers OK
- Simpler and faster

# Advantages of NoSQL

- Cheap, easy to implement (open source)

# Advantages of NoSQL

- Cheap, easy to implement (open source)
- Data Replication to multiple nodes

# Advantages of NoSQL

- Cheap, easy to implement (open source)
- Data Replication to multiple nodes
- No single point of failure

# Advantages of NoSQL

- Cheap, easy to implement (open source)
- Data Replication to multiple nodes
- No single point of failure
- Down nodes easily replaced



# Advantages of NoSQL

- Cheap, easy to implement (open source)
- Data Replication to multiple nodes
- No single point of failure
- Down nodes easily replaced
- Can scale up and down

# Advantages of NoSQL

- Cheap, easy to implement (open source)
- Data Replication to multiple nodes
- No single point of failure
- Down nodes easily replaced
- Can scale up and down
- No Schema

# Disadvantages of NoSQL

- JOIN Queries - Restructure Data ?

# Disadvantages of NoSQL

- JOIN Queries - Restructure Data ?
- Aggregate Queries - Pre-computation ?

# Disadvantages of NoSQL

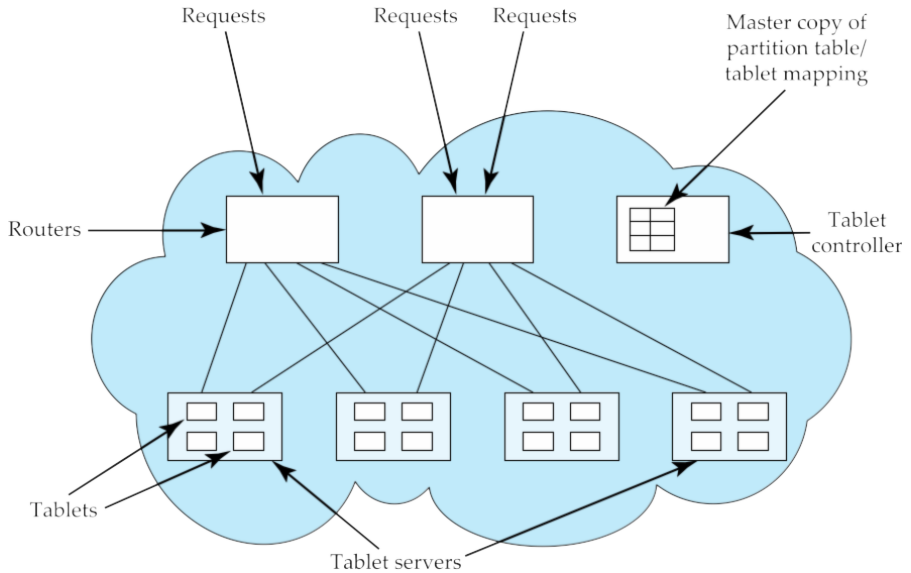
- JOIN Queries - Restructure Data ?
- Aggregate Queries - Pre-computation ?
- ACID Transaction - Available Partly.

# Disadvantages of NoSQL

- JOIN Queries - Restructure Data ?
- Aggregate Queries - Pre-computation ?
- ACID Transaction - Available Partly.
- SQL is unavailable - ORM and other query languages

# Disadvantages of NoSQL

- JOIN Queries - Restructure Data ?
- Aggregate Queries - Pre-computation ?
- ACID Transaction - Available Partly.
- SQL is unavailable - ORM and other query languages
- Change applications already written in SQL





# Basic Definitions and alternatives

- Entity - A row or Document or JSON Object

# Basic Definitions and alternatives

- Entity - A row or Document or JSON Object
- Entity Type - Table or collection or JSON Class

# Basic Definitions and alternatives

- Entity - A row or Document or JSON Object
- Entity Type - Table or collection or JSON Class
- Key-Value - Cell or Value

# Basic Definitions and alternatives

- Entity - A row or Document or JSON Object
- Entity Type - Table or collection or JSON Class
- Key-Value - Cell or Value
- Metadata - stores keys, time-stamp, id, etc

# Basic Definitions and alternatives

- Entity - A row or Document or JSON Object
- Entity Type - Table or collection or JSON Class
- Key-Value - Cell or Value
- Metadata - stores keys, time-stamp, id, etc
- Indexing - Sorted list of keys and addresses

# Basic Definitions and alternatives

- Entity - A row or Document or JSON Object
- Entity Type - Table or collection or JSON Class
- Key-Value - Cell or Value
- Metadata - stores keys, time-stamp, id, etc
- Indexing - Sorted list of keys and addresses
- Expando - Adding key-value to existing document

# Basic Definitions and alternatives

- Entity - A row or Document or JSON Object
- Entity Type - Table or collection or JSON Class
- Key-Value - Cell or Value
- Metadata - stores keys, time-stamp, id, etc
- Indexing - Sorted list of keys and addresses
- Expando - Adding key-value to existing document
- Polymodel - Inheritance of Entity Type

# Basic API

- `get(key)` – Extract the value given a key
- `put(key, value)` – Create or update the value given its key
- `delete(key)` – Remove the key and its associated value
- `execute(key, operation, parameters)` – Invoke an operation to the value (given its key)



# Other Databases

- Domain Specific Databases

# Other Databases

- Domain Specific Databases
- Relaxes further domain specific constraints

# Other Databases

- Domain Specific Databases
- Relaxes further domain specific constraints
- Graph Databases - Titan, Neo4j

# Other Databases

- Domain Specific Databases
- Relaxes further domain specific constraints
- Graph Databases - Titan, Neo4j
- Document Database - CouchDB, S3

# Other Databases

- Domain Specific Databases
- Relaxes further domain specific constraints
- Graph Databases - Titan, Neo4j
- Document Database - CouchDB, S3
- Hash-Like Databases - Voldemort, Scalaris

# Other Databases

- Domain Specific Databases
- Relaxes further domain specific constraints
- Graph Databases - Titan, Neo4j
- Document Database - CouchDB, S3
- Hash-Like Databases - Voldemort, Scalaris
- Key Value Systems like Lucene

# Tips and Tricks

- First ensure you need NoSQL. Petabyte scale websites like Quora are still using MySQL
- High Read / Update frequency 500 writes / sec / GB RAM
- Log Analysis, Creating Warehouses, Backup, etc
- Prefer Cloud Services to avoid maintenance
- Index everything properly
- Entity update close to 1 write / sec

# Questions ?

Questions ?



# Acknowledgments

- Thank You IEEE
- Thank You Karthik Senthil
- Thank You for attending

Ashish Kedia - 12IT14